



REPORT ASSIGNMENT 1

Analyzing NBA data

Emery Ong A0136591B

Emile Brès A0132365L

Simon Helmlinger A0134470M

Juan Manuel Muñoz Perez A0134739X

Monday 9th February, 2015

Contents

1	Part 1 - webscraping	3
1.1	Question 1	3
1.2	Question 2	3
1.2.1	Question 2.a	3
1.2.2	Question 2.b	3
1.2.3	Question 2.c	3
1.3	Question 3	4
1.3.1	Question 3.a	4
1.3.2	Question 3.b	4
1.3.3	Question 3.c	5
1.3.4	Question 3.d	5
1.4	Question 4	5
1.4.1	Question 4.a	5
1.4.2	Question 4.b	6
1.5	Question 5	6
2	Part 2	8
2.1	Question 1	8
2.2	Question 2	8
2.3	Question 3	9
2.3.1	Question 3.a	9
2.3.2	Question 3.b	9
2.3.3	Question 3.c	9
2.3.4	Question 3.d	9
2.4	Question 4	9
2.4.1	Question 4.a	10
2.4.2	Question 4.b	10
2.4.3	Question 4.c	10
2.5	Question 5	10

1 Part 1 - webscraping

1.1 Question 1

We want to retrieve data for the players and teams. The starting URL to do this are www.basketball-reference.com/players/ and www.basketball-reference.com/teams/ respectively.

1.2 Question 2

We decided to scrap the pages of all the players and not only the active ones, as the data could be useful to assess the salaries of the players. However, we also created a function `active_players_BS` which return which players are active and which are not.

The players are classified according to the first letter of their last name. We first retrieve all the letters which have players whose name starts with from <http://www.basketball-reference.com/players/>. For example, there is no player whose name starts with *x*. For a letter *c*, we then download the page <http://www.basketball-reference.com/players/c> and extract the player's page links.

1.2.1 Question 2.a

The regular expression we use to retrieve the letters is:

```
1 <a href="/players/([a-z]+)/">[A-Z]</a></td>
```

The regular expression we use to identify the links to the player's page is the following:

```
1 [^p]><a href="/players/./.+)"
```

1.2.2 Question 2.b

The code we use to retrieve the letters with BeautifulSoup4 is:

```
1 for row in soup('td', {'class': 'align_center bold_text valign_bottom
2                               xx_large_text'}):
3     letter = str(row.a.get('href').split('/')[2])
```

The code we use to identify the links with BeautifulSoup4 is the following:

```
1 for player in soup.tbody.find_all('tr'):
2     link = str(player.td.a.get('href'))
```

The whole code can be found at [Code/Part1/BS](#) and [Code/Part1/Regex](#) for the BeautifulSoup4 and Regex respectively.

1.2.3 Question 2.c

We chose to use the regex code because it is faster and use less memory. In fact, we ran 100 instances of both versions and obtained the following results:

Method	Time	Memory size
regex	1.167 s	2.6 MB
BS	3.911 s	3.7 MB

Those results were obtained using pre-downloaded pages to avoid measuring the download times stored in `Data/Part1/Metrics/Pages`.

Furthermore, the regex code is more readable since it contains less functions that need to be called.

1.3 Question 3

1.3.1 Question 3.a

Using the links parser from the code written in Part 1 Question 2, the entire player database was parsed using BeautifulSoup and basic profile information were scraped from the website. A standard template was designed to hold variables such as `playerid` (derived from the player's HTML link was used as an unique identifier), name, positions, shooting hand, height, weight, birth date, city, state/country of birth, experience and death date. When there are no values listed in the profile page, null values are assigned to the variables.

1.3.2 Question 3.b

The method used to scrap the basic statistics and salaries data is as follows:

- **Scraping phase** by analyzing the structure of the tables (HTML tags) into .csv files
- Adding an unique **Player ID** at each table constructed as follows: `/players/b/bryanko01.html` becomes `bbryanko01`. Each player then can be accessed with its unique key
- **Cleaning phase** where some column's formats are modified so that they can be loaded and manipulated in the SQLite database

Among all statistics of basketball players, scraping some tables is sufficient to do the analysis requested in Part 2. The tables selected as player's statistics are below:

- Totals
- Per Game

These tables deal with statistics by season (Totals) and by game (Per Game) which is enough to conduct the requested analysis. However, since tables are not exactly identical (HTML tags), scraping them wasn't as easy as expected. Numerous specialties have had to be taken in consideration.

For example, here is referenced a special case encountered in Totals and Per Game tables when the player qualifies to All-Star games. Next to the Season value, a star appears. Below the HTML code corresponding to the Totals table of Kobe Bryant.

```
1 <td align="left" >
2     <a href="/players/b/bryanko01/gamelog/1998/">1997-98</a>
3     <span class="bold_text" style="color:#c0c0c0">&nbsp;&#x2605;</span>
4 </td>
```

Scraping these web pages in which the code should catch all these specialties was tricky. For further information about the code providing the player's statistics scraper, see the `player_statistics_BS.py` file at `Code/BS/`.

The output `totals_final.csv` and `per_game_final.csv` tables can be find in `Data/Part1/` folder.

1.3.3 Question 3.c

Scraping the player's salaries tables require to scrap two main tables:

- Salaries (previous salaries)
- Contract (current and futures contracts if any)

Thus, the method is to merge these two tables. Here again, some specialties on tables structure render the scraping tricky. Indeed, Contract table doesn't have a `thead`, `tbody` and `tfoot` tag. So its scraping is done differently. Furthermore, salary's format is \$15,000,000 for instance. Unchanged, salary cannot be loaded to the SQLite database. Dollar symbol and commas have been removed so that the new salary's format becomes 15000000 (Cleaning phase).

The output `salaries_final.csv` table can be find in `Data/Part1/` folder.

1.3.4 Question 3.d

Question 3.A (see 1.3.1) can be repeated using Regex. A sample skeleton has been written (`profile_parser_regex.py` in `Code/Part1/Regex/`) to parse the player profile using Regex. Using BeautifulSoup4 to parse seems to be easier because we can make use of existing methods to call various functions like `soup.find`, `soup.findNext` and `soup.get_text`. Using these functions make the code more readable and easily understandable. It allows for easy debugging should the program break due to an uncaught exception. Regex tends to be more rigid and inflexible. BeautifulSoup4 allows us to traverse the parse tree to look for specific objects at specific part of the HTML document. With regex, we have to search the entire document for each search item.

We prefer to use BeautifulSoup4 to parse HTML and obtain the strings we are interested in. Then we can use regex to retrieve the information in those strings. We feel then that the best way is to use a combination of both.

1.4 Question 4

We use BeautifulSoup to scrap the franchise pages. More precisely, we use this package to navigate into the HTML tree, but we still need the regex package to make simple operations on string variables that can be found on the leaves of the tree.

1.4.1 Question 4.a

First of all, we want to scrap the basic team information. These are contained in the header of the franchise page. The HTML of this header can be found in the first `div` tag whose class is `mobile_text`. This tag is very badly organized because it consists in a succession of tags included in each other. Instead of having a list of sibling tags, all its successors are included in each other. Therefore, the BeautifulSoup code can be very tedious to write as each tag can only be reached by going down the hierarchy. Fortunately, we can avoid this problem with the function `get_text()` that automatically searches the text element in the HTML code and concatenate them. When we apply it to the header of the franchise page, we obtain a string containing all the basic information of the team. If we take the example of Atlanta Hawks franchise page, we get the following text :

Location Atlanta, Georgia

Team Names Atlanta Hawks, St. Louis Hawks, Milwaukee Hawks, Tri-Cities Blackhawks

Seasons 66; 1949-50 to 2014-15

Record 2584-2609, .498 W-L %

Playoff Appearances 43

Championships 1

Once we have this text, we just have to use the appropriate regex expressions to retrieve all the information we are interested in. In our project, we scrap them all, except the team names (we will scrap them in another table). You can notice that some pieces of information are not atomic. The **Location** field for example, contains two information: city and state. Thus, we split all information that are not atomic into several pieces. **Location** is split into city and state, **Seasons** into number of seasons, first season and last season and **Record** into number of wins, number of losses and win-loss percentage. We choose to identify a season with the civil year of the beginning of the season that is to say that the season 2014-15 is denoted as 2014. Besides, we have to choose a unique identifier for each franchise, since each one can have several names throughout the years. This id is a three letters abbreviation of the franchise name as it is used in the url of the pages of the franchise. For example, the abbreviation of Atlanta Hawks is **ATL** because the url of its page is <http://www.basketball-reference.com/teams/ATL/>.

All the basic information of all teams are scraped by the script named `team_scraper_BS` and they are stored in the same table named `teams_basic_info.csv`.

1.4.2 Question 4.b

Now we want to retrieve all the team statistics by season. These data are available on the same page as the team basic information and are already well-organized as we can find them in a table below the header. As this scraping job is very similar to what we did in 1.2, we will not describe it again. However, we can point out the fact that the name of the team can change throughout the seasons and that we don't use it as it is. Instead, we use a three letters abbreviation that can be found in the url pointing to the team roster of the corresponding season. The table containing the scraping result thus has both **Franchise_id** and **Team_id** columns. All statistics of all teams are scraped by the script gathered in a csv file named `teams_statistics`. These data are not sufficient to answer all questions of Part 2. In order to get players experience, we also have to scrap rosters of all teams for each season. These data are available in `roster_statistics.csv`.

1.5 Question 5

We thought that some of the highest salaries may be explained by a phenomenon of exaggerated enthusiasm for players who are playing well during the regular season but perform poorly once in the playoffs. To investigate this hypothesis, we decided to scrap the playoff statistics, that is to say statistics of teams that reached the final part of the competition. We thought that, by focusing on the best teams for each season, we would be able to assess the players who perform well when it really matters (in the playoffs). These data were scraped on the page <http://www.basketball-reference.com/playoffs/> with the same technique as in 1.2 and 1.4, since they are also stored in a well-structured table on the website.

Our recommendation to team owners is to consider players that appear in this page differently than other players. Performance during the playoffs should be more valued than performance during the regular season. Also the owners should pay attention to discrepancies between regular season production and playoffs production. It would help them not to pay salaries that are not deserved (to player performing worse during

the playoffs) and use this money for more efficient players (who perform well in the playoffs).

2 Part 2

2.1 Question 1

The database is composed of 7 tables which are given below. *Italic attributes are the keys of the table.*

Profile table *playerid* (varchar[12]), name (varchar[56]), position1 (varchar[2]), position2 (varchar[2]), position3 (varchar[2]), shoots (char[1]), height (integer), weight (integer), dob (date), city (varchar[56]), state (varchar[56]), experience (integer) and dod (date).

Totals table *playerid* (varchar[12]), season (integer), age (integer), team (char[3]), position (varchar[6]), g (integer), gs (integer), mp (integer), fg (integer), fga (integer), fg_percent (double), three_point (integer), three_point_a (integer), three_point_percent (double), two_point (integer), two_point_a (integer), two_point_percent (double), efg_percent (double), ft (integer), fta (integer), ft_percent (double), orb (integer), drb (integer), trb (integer), ast (integer), stl (integer), blk (integer), tov (integer), pf (integer) and pts (integer).

Per Game table The schema of this table is the same as Totals table except the *efg_percent* attribute which is absent from Per Game table.

Salary table *playerid* (varchar[12]), season (integer), team (varchar[24]), league (char[3]) and salary (integer).

Team Information table *franchiseid* (char[3]), city (varchar[24]), state (varchar[24]), season (integer), first_season (integer), last_season (integer), wins (integer), losses (integer), winlose_percent (double), playoff_app (integer) and championships (integer).

Team Map table *teamid* (varchar[12]) and team_name (varchar[128]).

Team Statistics table *franchiseid* (varchar[12] NOT NULL), *season* (integer), league (char[3]), *teamid* (char[3] NOT NULL), win (integer), loss (integer), wl_percent (double), finish (integer), srs (double), pace (double), rel_pace (double), ortg (double), rel_ortg (double), drtg (double), rel_drtg (double) and playoffs (varchar[128]).

Team Roster table *teamid* (char[3] NOT NULL), season (integer), number (varchar[6]), *playerid* (varchar[12] NOT NULL), position (varchar[6]) and experience (integer).

Active Players table *playerid* (varchar[12]) and active (varchar[12]).

2.2 Question 2

All the results below were found with the SQLite queries stored in the `Code/Part2/Q2` directory.

Active players We found 379 active players in 2011-2012 season among the 4288 players in total.

Distribution in each position The number of players playing in each position is as follow (we decide to consider players who played in two positions separately) :

Center 75

Center and Power Forward 1

Power Forward 82

Power Forward and Small Forward 1

Point Guard 69

Small Forward 78

Small Forward and Power Forward 1

Shooting Guard 78

Average age, weight and experience We found that the average age of players was 26.2 years, their average weight 220 lbs (= 99.2kg) and their average experience 4 years and 6 months.

Average salary in season and career The average salary of a player in 2011-2012 season was \$4,699,756.00 and its average career salary \$26,463,149.51.

2.3 Question 3

All the SQLite queries that allowed us to get the following results are available in the folder Code/Part2/Q3

2.3.1 Question 3.a

We interpreted this question as we needed to select a list of current active players who had a salary in 2011-2012 season, and list the top 10% from the reduced sample size. The results are available in the following file in : Data/Part2/Q3/question3a_best_paid_players.csv

2.3.2 Question 3.b

Bottom 10% worst paid players who are currently active and had a salary from 2011-2012 season are available in : Data/Part2/Q3/question3b_worst_paid_players.csv

2.3.3 Question 3.c

Middle 50% by pay of players who are currently active and had a salary from 2011-2012 season : Data/Part2/Q3/question3c_middle_paid_players.csv

2.3.4 Question 3.d

The money that was paid to all users by season, the number of active players by season and the average salary per player by season are available in the following csv file : Data/Part2/Q3/question3d_miscellaneous_salaries.csv

2.4 Question 4

The answers to the following questions are provided in .csv files in the ZIP folder. For each question, we will provide the name of the file containing the answer as well as the location of the source code in the ZIP folder.

2.4.1 Question 4.a

The SQLite queries are in `Code/Part2/Q4/sql_question_part2_Q4a.sql`. The answers are in `Data/Part2/Q4/question_4a_average_salaries.csv` and `Data/Part2/Q4/question_4a_variance_of_salaries.csv`.

2.4.2 Question 4.b

The SQLite queries are in `Code/Part2/Q4/sql_question_part2_Q4b.sql`. The answers are in `Data/Part2/Q4/question_4b_average_age.csv` and `Data/Part2/Q4/question_4b_average_and_variance_experience.csv`.

2.4.3 Question 4.c

The metrics above can be represented in a cross-tabulation format through the use of the case statement. The query is called for one team and then repeated for all the teams in 2011-2012 season and each result is concatenated using `UNION ALL`.

The query can be found in `Data/Part2/question_4c_cross_tabulation_format.csv`. The query shows the average age where team name is the first column for each row and year is the column. The query is then repeated and an `UNION ALL` is performed to join each query. The process is repeated for each team (ATL BOS CHA CHI CLE DAL DEN DET GSW HOU IND LAC LAL MEM MIA MIL MIN NJN NOH NYK OKC ORL PHI PHO POR SAC SAS TOR UTA WAS).

The results for 3 teams are shown in `Data/Part2/question_4c_cross_tabulation_format.csv`. For teams who changed names between 2002 and 2012 (for example Charlotte Hornets became New Orleans Hornets), both names would be in the table and `NULL` values would be inserted for the years when the team did not exist. The rest of the teams have not been done because it is very time consuming to write the code for each team.

2.5 Question 5

We could use the in-game statistics for each player to assess his salary (points, rebounds, assists...). However, these statistics do not always reflect properly the impact of a player. Since basketball is a team sport, the performance of a player relies heavily on the performance of his teammates. Advanced statistics have been developed to address these issues.

In particular, `win shares` are an estimation of the number of wins contributed by a player over a season. For instance, LeBron James had the most `win shares` between 2008 and 2013 and won 4 MVP titles in these 5 years. James' worth is well-known among basketball circles but there are some players who are less recognized and have good `win shares` values.

Furthermore, the `win shares` statistics are already on `www.basketball-reference.com` for each player in the table `Advanced`. Retrieving that data would only require adding a line to `get_statistics_player` function in `Code/Part1/BS/player_statistics.py`.

We would thus advise team owners to look for discrepancies in the correlation between salary and `win shares`.